

STA 380 Homework 1

Nicole Erich, Anying Li, Daniel Peng & Rachel Wang

August 5, 2016

Probability practice

Part A.

Based on law of total probability: $P(Y) = P(Y|RC) * P(RC) + P(Y|TC) * P(TC)$

Here: $P(Y) = 0.65$ $P(Y|RC) = 0.5$ $P(RC) = 0.3$ $P(TC) = 0.7$

Therefore: $P(Y|TC) = 0.7142857$

So about 71.43% of people who are truthful clickers answered yes.

Part B.

Based on Bayes' Rule: $P(A|B) = [P(A) * P(B|A)] / P(B)$

Here:

Event A is someone has the disease; event B is someone's test result is positive. We want to know $P(A|B)$.

$P(A) = 0.000025$ $P(B|A) = 0.993$ We can calculate $P(B)$ based on law of total probability: $P(B) = P(B|A) * P(A) + P(B|\text{not } A) * P(\text{not } A) = 0.993 * 0.000025 + (1 - 0.9999) * (1 - 0.000025) = 0.0001248225$

So $P(A|B) = (0.000025 * 0.993) / 0.0001248225 = 0.1988824$

Q1 Exploratory analysis: green buildings

```
all.buildings = read.csv('https://raw.githubusercontent.com/jgscott/STA380/master/data/greenbuildings.csv')
green.buildings = subset(all.buildings, all.buildings$EnergyStar == 1 | all.buildings$LEED == 1)
not.green = all.buildings[!(all.buildings$CS_PropertyID %in% green.buildings$CS_PropertyID),]
```

To estimate the economic impact of a green certificate, we had to calculate the expected extra profit brought in with it. To do so, We needed to find out the additional cost and revenue associated with a green building. The extra cost for this property is the \$5 Million premium (\$100M x 5%). The extra revenue per year would be additional rent/sqft-year x size of the building (250,000 sqft). In different clusters, we might value a green certificate differently. Therefore, we could not simply find the median rent for regular buildings and green buildings and subtract one from the other. Naturally, we were going to find the difference between two types of buildings by clusters. In this case, we treated all the regular buildings in a certain cluster as control group so we could see the effect of a certificate.

```
#calculate the average rent within each cluster
notgreen.mean.rent = aggregate(not.green$Rent, list(not.green$cluster), mean)
green.mean.rent = aggregate(green.buildings$Rent, list(green.buildings$cluster), mean)

#remove the clusters without any green buildings
```

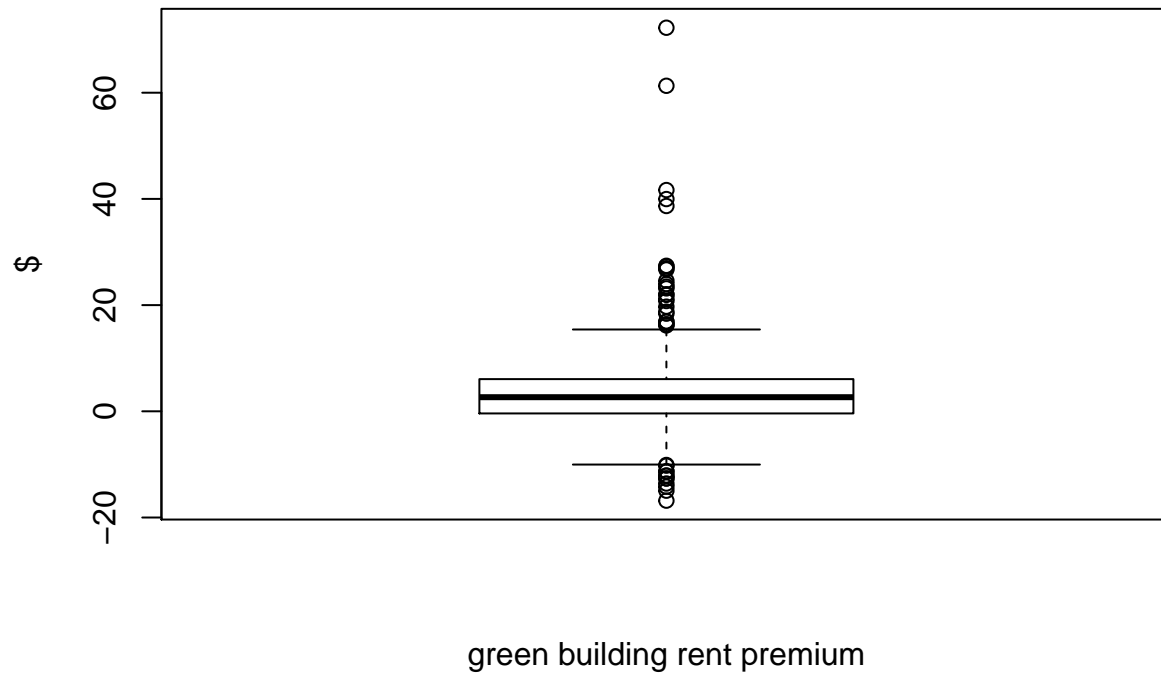
```

notgreen.mean.rent = notgreen.mean.rent[which(notgreen.mean.rent[,1] %in% green.buildings$cluster),]

#remove the clusters without any regular buildings
green.mean.rent = green.mean.rent[which(green.mean.rent[,1] %in% notgreen.mean.rent[,1]),]

rent.diff = green.mean.rent - notgreen.mean.rent
boxplot(rent.diff$x, ylab = '$', xlab = 'green building rent premium')

```



```

rent = median(rent.diff$x)
c('expected additional rent:', rent)

```

```
## [1] "expected additional rent:" "2.66"
```

From the boxplot above, we can see the majority of the rent premium is concentrated around \$2, with many outliers. So It's better to use median of \$2.66 for estimating the new building.

We used the same approach to calculate the expected occupancy rate change for the new property.

```


#calculate the average leasing rate within each cluster
notgreen.mean.lr = aggregate(not.green$leasing_rate, list(not.green$cluster), mean)
green.mean.lr = aggregate(green.buildings$leasing_rate, list(green.buildings$cluster), mean)

#remove the clusters without any green buildings
notgreen.mean.lr = notgreen.mean.lr[which(notgreen.mean.lr[,1] %in% green.buildings$cluster),]

#remove the clusters without any green buildings
green.mean.lr = green.mean.lr[which(green.mean.lr[,1] %in% notgreen.mean.lr[,1]),]

lr.diff = green.mean.lr - notgreen.mean.lr
boxplot(lr.diff$x, ylab = 'percentage', xlab = 'leasing rate difference between green vs non-green build

```



leasing rate difference between green vs non-green building


```
lr = median(lr.diff$x)
c('expected additional occupancy:', lr)
```

```
## [1] "expected additional occupancy:" "8.08"
```

Interestingly, the distribution of the occupancy difference is very concentrated around 10. It seems people prefer green buildings.

In order to find a good estimate for the occupancy rate for the new building, we plotted the distribution of all existing green buildings. The median of 92.93 looked like a good choice based on the boxplot.

```
boxplot(green.mean.lr[,2], ylab = 'percentage', xlab = 'green building leasing rate')
```



green building leasing rate

In conclusion, our finding is very similar to the original analysis even though with different approach: rent premium = \$2.66, occupancy rate = 92.93%.

```
c('Number of Years to Break Even: ', 500000 / (2.66 * 250000 * .9293))
```

```
## [1] "Number of Years to Break Even: " "8.0908178117736"
```

```
c('Additional Annual Revenue for Green Certification:', (2.66 * 250000 * .9293))
```

```
## [1] "Additional Annual Revenue for Green Certification:"
```

```
## [2] "617984.5"
```

Q2 Bootstrapping

et up and create the function for calculating percent returns.

```
rm(list=ls())
library(mosaic)
```

```
## Loading required package: dplyr
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

## Loading required package: lattice

## Loading required package: ggplot2

## Loading required package: mosaicData

## Loading required package: Matrix

##
## The 'mosaic' package masks several functions from core packages in order to add additional features.
## The original behavior of these functions should not be affected by this.

##
## Attaching package: 'mosaic'

## The following object is masked from 'package:Matrix':
##
##   mean

## The following objects are masked from 'package:dplyr':
##
##   count, do, tally

## The following objects are masked from 'package:stats':
##
##   binom.test, cor, cov, D, fivenum, IQR, median, prop.test,
##   quantile, sd, t.test, var

## The following objects are masked from 'package:base':
##
##   max, mean, min, prod, range, sample, sum
```

```
library(fImport)
```

```
## Loading required package: timeDate
```

```
## Loading required package: timeSeries
```

```
library(foreach)
my_favorite_seed = 1234567
set.seed(my_favorite_seed)
YahooPricesToReturns = function(series) {
  mycols = grep('Adj.Close', colnames(series))
  closingprice = series[,mycols]
  N = nrow(closingprice)
  percentreturn = as.data.frame(closingprice[2:N,]) / as.data.frame(closingprice[1:(N-1),]) - 1
```

```

mynames = strsplit(colnames(percentreturn), '.', fixed=TRUE)
mynames = lapply(mynames, function(x) return(paste0(x[1], ".PctReturn")))
colnames(percentreturn) = mynames
as.matrix(na.omit(percentreturn))
}

```

Import the 5 asset classes and calculate their respective standard deviations.

```

mystocks = c("SPY", "TLT", "LQD", "EEM", "VNQ")
myprices = yahooSeries(mystocks, from='2011-01-01', to='2016-08-03')
myreturns = YahooPricesToReturns(myprices)
sigma_SPY = sd(myreturns[,1])
sigma_SPY

```

```
## [1] 0.009650712
```

```

sigma_TLT = sd(myreturns[,2])
sigma_TLT

```

```
## [1] 0.009385556
```

```

sigma_LQD = sd(myreturns[,3])
sigma_LQD

```

```
## [1] 0.003428545
```

```

sigma_EEM = sd(myreturns[,4])
sigma_EEM

```

```
## [1] 0.01411309
```

```

sigma_VNQ = sd(myreturns[,5])
sigma_VNQ

```

```
## [1] 0.01136797
```

The standard deviations show that Emerging-market equities (EEM) and Real estate (VNQ) are the most volatile asset classes. We construct an aggressive portfolio with a 50/50 split between EEM and VNQ, a safe portfolio with the other three asset classes, and an even split by distributing 20% of money in each of the five ETFs.

```

weights_safe = c(0.3, 0.3, 0.4, 0.0, 0.0)
weights_even = c(0.2, 0.2, 0.2, 0.2, 0.2)
weights_aggressive = c(0.0, 0.0, 0.0, 0.5, 0.5)

```

Use bootstrap resampling to estimate 4-week VaR of each of the three portfolios at the 5% level.

```

n_days = 20
wealth_tracker_safe = rep(0, 5000)
wealth_tracker_even = rep(0, 5000)
wealth_tracker_aggressive = rep(0, 5000)
sim1 = foreach(i=1:5000, .combine='rbind') %do% {
  totalwealth = 100000
  holdings_safe = weights_safe * totalwealth
  holdings_even = weights_even * totalwealth
  holdings_aggressive = weights_aggressive * totalwealth
  for(today in 1:n_days) {
    return.today = resample(myreturns, 1, orig.ids=FALSE)
    holdings_safe = holdings_safe + holdings_safe*return.today
    holdings_even = holdings_even + holdings_even*return.today
    holdings_aggressive = holdings_aggressive + holdings_aggressive*return.today
    totalwealth_safe = sum(holdings_safe)
    totalwealth_even = sum(holdings_even)
    totalwealth_aggressive = sum(holdings_aggressive)
    holdings_safe = weights_safe * totalwealth_safe
    holdings_even = weights_even * totalwealth
    holdings_aggressive = weights_aggressive * totalwealth_aggressive
  }
  wealth_tracker_safe[i]=totalwealth_safe
  wealth_tracker_even[i]=totalwealth_even
  wealth_tracker_aggressive[i]=totalwealth_aggressive
}

```

```

var_safe = quantile(wealth_tracker_safe, 0.05) - 100000
var_safe

```

```

##          5%
## -2045.256

```

```

var_even = quantile(wealth_tracker_even, 0.05) - 100000
var_even

```

```

##          5%
## -931.2584

```

```

var_aggressive = quantile(wealth_tracker_aggressive, 0.05) - 100000
var_aggressive

```

```

##          5%
## -7828.275

```

The VaR analysis tells us that the theoretically safer portfolio is in fact more volatile than the even split. One reasonable explanation is that the risk of a portfolio is decided by its overall diversity rather than the standalone volatility of its components. Therefore, event split (best diversity) outperforms the “safer” portfolio comprised by the “safer” asset classes.

Q3 Market segmentation

Step 1. In terms of data processing, first we get the frequency of each category by dividing each data by corresponding row sum, so we can get the relative weight of each category in each twitter. Then we try to limit the noise from categories such as “chatter”, “spam”, “adult” and “uncategorized” because these are not really user twitter with helpful information. We decide to get rid of the records if the sum of the frequency of the four noise categories is bigger than 0.5. In the end, we scale and centralize the data to get the data prepared for clustering.

```
library(cluster)
library(fpc)
library(flexclust)
```

```
## Loading required package: grid
```

```
## Loading required package: modeltools
```

```
## Loading required package: stats4
```

```
library(foreach)
library(ggplot2)
```

```
social_marketing <- read.csv("https://raw.githubusercontent.com/jgscott/STA380/master/data/social_marke
```

```
# Center/scale the data
X_freq = social_marketing/rowSums(social_marketing)
dim(X_freq)
```

```
## [1] 7882    36
```

```
any(is.na(X_freq))
```

```
## [1] FALSE
```

```
#take the record out if frequency of uninformative category is over 0.5
X_freq['sum'] = X_freq$chatter +X_freq$spam +X_freq$adult +X_freq$uncategorized
X_freq = subset(X_freq, X_freq$sum <= 0.5)
dim(X_freq)
```

```
## [1] 7849    37
```

```
X_freq <- X_freq[, -37]
```

```
social_marketing_scaled = scale(X_freq, center=TRUE, scale=TRUE)
```

Step 2. We decide to begin with a k-mean clustering to see whether we can find something interesting. To use k-mean clustering, finding the appropriate k is our first task. To get the optimal k, we tried two approaches. First approach is to calculate the within group sum of square(wss). It's obvious that when k increases, wss will keep decreasing. But after plotting the relationship between k and wss, we can see that with 6 clusters the wss has already decreased significantly. As a second approach, we also calculated CH index that we


discussed at class to decide the optimal k . For the seed we set, we get a optimal k equal to 12. But we should also notice that after 6 clusters, CH index does not improve significantly as k keeps increasing. Meanwhile, too many clusters can make it hard to interpret the meaning behind the clustering. Therefore, we decide to use $k = 6$ for the clustering.

```
#decide optimal k based on within group sum of square
k.max <- 15 # Maximal number of clusters
data <- social_marketing_scaled

# Compute and plot wss for k = 2 to k = 15
set.seed(1234567)
wss <- sapply(1:k.max,
              function(k){kmeans(data, k, nstart=50)$tot.withinss})
```

```
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
```

```
plot(1:k.max, wss,
     type="b", pch = 19, frame = FALSE,
     xlab="Number of clusters K",
     ylab="Total within-clusters sum of squares")
abline(v = 6, lty = 2)
```



#6 clusters are suggested

#decide optimal k based on within CH index/Average silhouette method

```
sil <- rep(0, k.max)
```

Compute the average silhouette width for k = 2 to k = 15

```
for(i in 2:k.max){
  set.seed(1234567)
  km.res <- kmeans(data, centers = i, nstart = 50)
  ss <- silhouette(km.res$cluster, dist(data))
  sil[i] <- mean(ss[, 3])
}
```

```
## Warning: did not converge in 10 iterations
```

```
## Warning: did not converge in 10 iterations
```

```
## Warning: did not converge in 10 iterations
```

```
## Warning: did not converge in 10 iterations
```

```
## Warning: did not converge in 10 iterations
```

```
## Warning: did not converge in 10 iterations
```

```
## Warning: did not converge in 10 iterations
```


```
## Warning: did not converge in 10 iterations
```

```
## Warning: did not converge in 10 iterations
```

```
## Warning: did not converge in 10 iterations
```

```
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
```

```
# Plot the average silhouette width
plot(1:k.max, sil, type = "b", pch = 19,
     frame = FALSE, xlab = "Number of clusters k")
abline(v = which.max(sil), lty = 2)
abline(v = 6, lty = 2)
```



```
# 12 clusters are suggested
```

Step3: After deciding the optimal k, we tried to build the clustering with both kmean and kmean++ for initialization, and it turns out that the error from these two methods are very close. It might be related with the seed we choose, and as we set nstart = 50 for kmean, it also helps improve the accuracy of kmean clustering.

After looking through the centers of the 6 clusters, we find:

second cluster has high positive weight on sports_fandom, food, family, religion, parenting and school, so this group may include married people who pay more attention to their family and parenting related topic.

third cluster has high positive weight on cooking, beauty and fashion, so this group should consist of younger woman and younger housewives who cook a lot and pay a lot of attention to beauty and fashion.

fourth cluster has high positive weight on politics, news, travel, computer and automotive, so this group might be younger man who are interested in politics, read lots of news online, loves computer and automotive and travels a lot.

fifth cluster has high positive weight on health_nutrition, outdoors and personal_fitness, so this group of people really care about nutrition and fitness, and they have lots of outdoor activity to keep fit.

sixth cluster has high positive weight on online_gaming, college_uni and sports_playing, this group is very likely to be college student whose main entertainment is online_gaming and sports.

first cluster's highest positive weight is on chatter, and nothing else really stands out. So we think this cluster may be just everything that is left out and is hard to get into any of the other market segment.

```
# fit cluster model with 6 centers
set.seed(1234567)
cluster_6 <- kmeans(data, centers=6, nstart=50)
names(cluster_6)
```

```
## [1] "cluster"      "centers"      "totss"       "withinss"
## [5] "tot.withinss" "betweenss"    "size"        "iter"
## [9] "ifault"
```

```
cluster_6$centers
```

```
##      chatter current_events      travel photo_sharing uncategorized
## 1 -0.3837312   -0.2108030 -0.18108941   -0.3361177   -0.05627069
## 2 -0.4339944   -0.2137770 -0.27376862   -0.4081966   -0.11315049
## 3 -0.4362083   -0.1669467 -0.28527970   -0.4193508   -0.21586985
## 4 -0.4570848   -0.1835754 -0.24617030    0.4358365    0.02503411
## 5 -0.3667735   -0.1015367  0.96771010   -0.4778611   -0.14669241
## 6  0.8083460    0.3361859 -0.04272988    0.5508299    0.22049860
##      tv_film sports_fandom      politics      food      family
## 1  0.19171228   -0.2934869 -0.3490226 -0.2142310 -0.09722020
## 2 -0.24938308   -0.3491980 -0.3302150  0.1256574 -0.25737526
## 3 -0.19392881    1.3885068 -0.3803047  1.2437855  0.81916145
## 4 -0.26599694   -0.4041387 -0.3604385 -0.4221893 -0.23678827
## 5 -0.08841854    0.1186687  1.7171596 -0.2293812 -0.09206439
## 6  0.27848491   -0.2577476 -0.2242795 -0.3011719 -0.07515292
##      home_and_garden      music      news online_gaming      shopping
## 1   -0.10447006  0.1318158 -0.3117212    2.3187349 -0.3735883
## 2   -0.09614825 -0.1361084 -0.1920087   -0.2121673 -0.2872899
## 3   -0.07896679 -0.1159392 -0.2653374   -0.2217606 -0.3069829
## 4   -0.12161566  0.1260787 -0.2868907   -0.2006455 -0.1782972
## 5   -0.05648961 -0.1676258  1.5227455   -0.2493419 -0.3689054
## 6    0.17212577  0.1184911 -0.2900471   -0.2418014  0.5919211
##      health_nutrition college_uni sports_playing      cooking      eco
## 1   -0.3685201    2.4361909    1.19273070 -0.3423883 -0.19162232
## 2    1.7709001   -0.3348953   -0.15758760  0.2118675  0.11849306
## 3   -0.3518720   -0.2973465   -0.15077221 -0.3471649 -0.04517358
## 4   -0.3035556   -0.2314739   -0.05938603  2.1879569 -0.20040270
## 5   -0.4049666   -0.2690893   -0.15882567 -0.3973473 -0.16046545
## 6   -0.3858046   -0.1603625   -0.08417465 -0.3663041  0.14143868
##      computers      business      outdoors      crafts      automotive      art
## 1 -0.20964361 -0.13506862 -0.24498869 -0.1739022 -0.18221011  0.009683922
## 2 -0.18708480 -0.15157659  1.08848837 -0.1213240 -0.28218207 -0.125017628
## 3 -0.09822232 -0.09561711 -0.25615098  0.2371491 -0.09281254 -0.072296995
## 4 -0.16071024 -0.03464582 -0.18582091 -0.1735611 -0.21043650 -0.090258403
## 5  0.76748918 -0.01639691 -0.01674134 -0.1374715  0.88205766 -0.189787115
## 6 -0.09624176  0.17122869 -0.31786781  0.1217041 -0.09250004  0.202846063
##      religion      beauty      parenting      dating      school
```

```
## 1 -0.2604777 -0.29058793 -0.3046890 -0.13145477 -0.33414491
## 2 -0.2926531 -0.28638240 -0.2667258  0.01814441 -0.29427076
## 3  1.6900537  0.06817419  1.5442861 -0.16015488  1.01732114
## 4 -0.3031804  1.76497590 -0.2809682 -0.13030015 -0.09246893
## 5 -0.2271477 -0.28854443 -0.1105381 -0.02522172 -0.20215764
## 6 -0.3011191 -0.21932366 -0.2993223  0.14411013 -0.07279344
##  personal_fitness  fashion small_business      spam      adult
## 1      -0.3646911 -0.2549782   -0.001005938  0.02812103 -0.06435910
## 2       1.6434108 -0.2610103   -0.183707292 -0.03294390 -0.07653893
## 3      -0.3149193 -0.1759703   -0.075352035 -0.04088964 -0.04968130
## 4      -0.3145744  2.0056300   -0.063076878 -0.04318552 -0.08976994
## 5      -0.3886946 -0.3366833   -0.055735098 -0.02648512 -0.08902164
## 6      -0.3411691 -0.1913383    0.169559661  0.05146539  0.14371580
```

```
head(sort(cluster_6$centers[1,], decreasing=TRUE), 10)
```

```
## college_uni online_gaming sports_playing      tv_film      music
## 2.436190862  2.318734947  1.192730698  0.191712282  0.131815757
##      spam      art small_business uncategorized      adult
## 0.028121027  0.009683922  -0.001005938  -0.056270690  -0.064359098
```


```
head(sort(cluster_6$centers[2,], decreasing=TRUE), 10)
```

```
## health_nutrition personal_fitness      outdoors      cooking
## 1.77090014      1.64341084      1.08848837  0.21186751
##      food      eco      dating      spam
## 0.12565742      0.11849306  0.01814441  -0.03294390
##      adult home_and_garden
## -0.07653893      -0.09614825
```

```
head(sort(cluster_6$centers[3,], decreasing=TRUE), 10)
```

```
## religion      parenting sports_fandom      food      school
## 1.69005373  1.54428612  1.38850676  1.24378546  1.01732114
## family      crafts      beauty      spam      eco
## 0.81916145  0.23714909  0.06817419  -0.04088964  -0.04517358
```

```
plotcluster(data, cluster_6$cluster)
```



```
summary(cluster_6)
```

```
##           Length Class  Mode
## cluster      7849  -none- numeric
## centers        216  -none- numeric
## totss           1  -none- numeric
## withinss        6  -none- numeric
## tot.withinss    1  -none- numeric
## betweenss       1  -none- numeric
## size            6  -none- numeric
## iter            1  -none- numeric
## ifault          1  -none- numeric
```

```
cluster_6$tot.withinss
```

```
## [1] 216511
```

```
# use kmean++ for clustering initialization
```

```
set.seed(123)
```

```
cluster_kmeansPP = cclust(data, k=6, control=list(initcent="kmeanspp"))
```

```
## Found more than one class "kcca" in cache; using the first, from namespace 'kernlab'
```

```
## Found more than one class "kcca" in cache; using the first, from namespace 'kernlab'
```

```
parameters(cluster_kmeansPP)
```

```
## Found more than one class "kcca" in cache; using the first, from namespace 'kernlab'
```

```

##      chatter current_events      travel photo_sharing uncategorized
## [1,] -0.3501608  -0.01499684 -0.05722007  -0.3558094   0.16234557
## [2,] -0.4409824  -0.14597170 -0.28299482  -0.4244248  -0.19939911
## [3,]  1.0741199   0.37351013 -0.08421317   0.8918532   0.13148004
## [4,] -0.3817482  -0.08836128  0.96096527  -0.4879916  -0.14753243
## [5,]  0.3838886  -0.10991406 -0.17527064  -0.3454694   0.26606373
## [6,] -0.4729166  -0.21541104 -0.25799758  -0.1132325  -0.07257773
##      tv_film sports_fandom      politics      food      family
## [1,]  0.95225957  -0.2634359 -0.3126114 -0.13556094 -0.174417047
## [2,] -0.18891548   1.3944322 -0.3803793  1.24708073  0.812576704
## [3,] -0.10148613  -0.2478624 -0.1868494 -0.36733918 -0.001533922
## [4,] -0.07911577   0.1188571  1.6848121 -0.24537197 -0.085534781
## [5,] -0.19515933  -0.3637056 -0.2798609 -0.34060771 -0.251642013
## [6,] -0.26162064  -0.3813797 -0.3598248 -0.04455022 -0.256628146
##      home_and_garden      music      news online_gaming      shopping
## [1,]  0.04173941  0.37446597 -0.2736503  1.3217010 -0.3234939
## [2,] -0.08066413 -0.10908600 -0.2653342  -0.2271706 -0.3141034
## [3,]  0.12156932  0.03263285 -0.3136961  -0.2503472  0.8939454
## [4,] -0.04238565 -0.17026960  1.5244505  -0.2417953 -0.3764808
## [5,]  0.11140532 -0.21388475 -0.2310183  -0.1578970 -0.3291675
## [6,] -0.09492986 -0.05242260 -0.2331256  -0.2143617 -0.2627523
##      health_nutrition college_uni sports_playing      cooking      eco
## [1,] -0.4078003  1.6272346  0.79360485 -0.3171857 -0.15801379
## [2,] -0.3491096  -0.3056514  -0.15349342 -0.3385289 -0.04634232
## [3,] -0.3791366  -0.2441802  -0.13523739 -0.3294268  0.21217071
## [4,] -0.3711231  -0.2777735  -0.17280931 -0.3874233 -0.16891296
## [5,] -0.2941602  -0.2247665  0.05124328 -0.3007900 -0.11737999
## [6,]  1.0729447  -0.3203662  -0.14243433  0.9758396  0.02158752
##      computers      business      outdoors      crafts      automotive
## [1,] -0.23835731  0.02309149 -0.260998947  0.02463569 -0.255819524
## [2,] -0.09436667 -0.09098105 -0.260533678  0.24027580 -0.089834426
## [3,] -0.05716320  0.15492124 -0.351864036  0.05032262 -0.004196339
## [4,]  0.76788984 -0.02360072  0.007706184 -0.14291931  0.894689517
## [5,] -0.08931371  0.15480875 -0.135262301  0.03377071 -0.271780691
## [6,] -0.18712317 -0.13012067  0.664835142 -0.11979289 -0.276542952
##      art      religion      beauty      parenting      dating      school
## [1,]  0.69370732 -0.2115014 -0.20603227 -0.3476384 -0.1898712 -0.2849086
## [2,] -0.07689029  1.6801603  0.07499968  1.5377590 -0.2149842  1.0184294
## [3,] -0.14106991 -0.3456032 -0.22536283 -0.2848370 -0.2350152 -0.1566664
## [4,] -0.17181440 -0.2298729 -0.29244394 -0.1170041 -0.0967113 -0.2163188
## [5,] -0.09123341 -0.1806937  0.15591244 -0.1380454  3.8320173  0.7329884
## [6,] -0.09865064 -0.2980649  0.44765725 -0.2749978 -0.1144050 -0.2382919
##      personal_fitness      fashion small_business      spam      adult
## [1,] -0.3813658 -0.1912280  0.22164207  0.03061494 -0.04254122
## [2,] -0.3099154 -0.1813723  -0.07825804 -0.03197029 -0.01629351
## [3,] -0.3262126 -0.2239218  0.08251367  0.05149493  0.08281286
## [4,] -0.3592700 -0.3488453  -0.07836819 -0.02069886 -0.03580816
## [5,] -0.2983826  0.5948847  0.08300706 -0.04883438  0.02420529
## [6,]  0.9767734  0.5398863  -0.13462788 -0.03134730 -0.03272741

```

```
cluster_kmeansPP@clusinfo
```

```

##      size av_dist max_dist separation
## 1 1161 5.581387 33.70702  3.506096

```

```
## 2 1120 4.828279 13.39611 3.498190
## 3 2029 5.165092 33.29725 3.096357
## 4 1181 5.151554 14.40995 3.496701
## 5 335 4.772608 10.41878 3.802397
## 6 2023 4.839438 17.06977 3.186441
```

```
print(apply(parameters(cluster_kmeansPP), 1, function(x) colnames(data)[order(x, decreasing=TRUE)[1:10]]))
```

```
##      [,1]      [,2]      [,3]      [,4]
## [1,] "college_uni" "religion" "chatter" "politics"
## [2,] "online_gaming" "parenting" "shopping" "news"
## [3,] "tv_film" "sports_fandom" "photo_sharing" "travel"
## [4,] "sports_playing" "food" "current_events" "automotive"
## [5,] "art" "school" "eco" "computers"
## [6,] "music" "family" "business" "sports_fandom"
## [7,] "small_business" "crafts" "uncategorized" "outdoors"
## [8,] "uncategorized" "beauty" "home_and_garden" "spam"
## [9,] "home_and_garden" "adult" "adult" "business"
## [10,] "spam" "spam" "small_business" "adult"
##      [,5]      [,6]
## [1,] "dating" "health_nutrition"
## [2,] "school" "personal_fitness"
## [3,] "fashion" "cooking"
## [4,] "chatter" "outdoors"
## [5,] "uncategorized" "fashion"
## [6,] "beauty" "beauty"
## [7,] "business" "eco"
## [8,] "home_and_garden" "spam"
## [9,] "small_business" "adult"
## [10,] "sports_playing" "food"
```

```
# Roll our own function
centers = parameters(cluster_kmeansPP)
kpp_residualss = foreach(i=1:nrow(data), .combine='c') %do% {
  x = data[i,]
  a = cluster_kmeansPP@cluster[i]
  m = centers[a,]
  sum((x-m)^2)
}
sum(kpp_residualss)
```

```
## [1] 222943.2
```

step4: We also tried Principal Component Analysis on the dataset. However, we find that the variance of the dataset cannot be simply explained by the top2 or top3 factors. In fact, top10 pc all have a pretty strong explanation power for the dataset variance. And after we print out the top words associated with first and second PC, it's harder to interpret the “market segment” compared with using just the clustering. Therefore we just use PCA as a supporting evidence to what we find in kmean clustering.

```
# PCA
pc = prcomp(X_freq, scale=TRUE)
```




```
set.seed(1234567)
loadings = pc$rotation
scores = pc$x
```


```
summary(pc)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation  1.70154 1.62129 1.54469 1.46176 1.41187 1.27794
## Proportion of Variance 0.08042 0.07302 0.06628 0.05935 0.05537 0.04537
## Cumulative Proportion 0.08042 0.15344 0.21972 0.27907 0.33444 0.37981
##              PC7      PC8      PC9     PC10     PC11     PC12
## Standard deviation  1.20439 1.14929 1.0916 1.05686 1.02932 0.99744
## Proportion of Variance 0.04029 0.03669 0.0331 0.03103 0.02943 0.02764
## Cumulative Proportion 0.42010 0.45679 0.4899 0.52092 0.55035 0.57799
##              PC13     PC14     PC15     PC16     PC17     PC18
## Standard deviation  0.986 0.98301 0.97282 0.95188 0.9467 0.92384
## Proportion of Variance 0.027 0.02684 0.02629 0.02517 0.0249 0.02371
## Cumulative Proportion 0.605 0.63183 0.65812 0.68329 0.7082 0.73189
##              PC19     PC20     PC21     PC22     PC23     PC24
## Standard deviation  0.89358 0.86179 0.84943 0.82782 0.8204 0.80437
## Proportion of Variance 0.02218 0.02063 0.02004 0.01904 0.0187 0.01797
## Cumulative Proportion 0.75407 0.77470 0.79475 0.81378 0.8325 0.85045
##              PC25     PC26     PC27     PC28     PC29     PC30
## Standard deviation  0.78796 0.77161 0.76796 0.75867 0.74529 0.73110
## Proportion of Variance 0.01725 0.01654 0.01638 0.01599 0.01543 0.01485
## Cumulative Proportion 0.86770 0.88424 0.90062 0.91661 0.93204 0.94688
##              PC31     PC32     PC33     PC34     PC35     PC36
## Standard deviation  0.69765 0.64682 0.61895 0.56561 0.55144 4.714e-15
## Proportion of Variance 0.01352 0.01162 0.01064 0.00889 0.00845 0.000e+00
## Cumulative Proportion 0.96040 0.97203 0.98267 0.99155 1.00000 1.000e+00
```


```
plot(pc)
```



```
biplot(pc)
```



```
qplot(scores[,1], scores[,2], xlab='Component 1', ylab='Component 2')
```



The top words associated with each component

```
o1 = order(loadings[,1])
colnames(X_freq)[head(o1,25)]
```

```
## [1] "religion"      "sports_fandom"  "parenting"
## [4] "food"          "school"         "family"
## [7] "news"          "automotive"     "crafts"
## [10] "politics"      "adult"          "computers"
## [13] "spam"          "art"            "travel"
## [16] "outdoors"      "dating"         "home_and_garden"
## [19] "small_business" "eco"            "tv_film"
## [22] "business"      "music"          "beauty"
## [25] "sports_playing"
```

```
colnames(X_freq)[tail(o1,25)]
```

```
## [1] "computers"      "spam"           "art"
## [4] "travel"         "outdoors"       "dating"
## [7] "home_and_garden" "small_business" "eco"
## [10] "tv_film"        "business"       "music"
## [13] "beauty"         "sports_playing" "current_events"
## [16] "personal_fitness" "health_nutrition" "online_gaming"
## [19] "uncategorized"  "college_uni"    "fashion"
## [22] "cooking"        "shopping"       "chatter"
## [25] "photo_sharing"
```

```
o2 = order(loadings[,2])
colnames(X_freq)[head(o2,25)]
```

```
## [1] "chatter"      "politics"      "travel"
## [4] "shopping"     "automotive"    "current_events"
## [7] "photo_sharing" "news"          "computers"
## [10] "tv_film"      "college_uni"   "small_business"
## [13] "business"     "online_gaming" "family"
## [16] "sports_playing" "home_and_garden" "sports_fandom"
## [19] "art"          "uncategorized" "music"
## [22] "crafts"       "school"        "spam"
## [25] "adult"
```

```
colnames(X_freq)[tail(o2,25)]
```

```
## [1] "small_business" "business"      "online_gaming"
## [4] "family"         "sports_playing" "home_and_garden"
## [7] "sports_fandom"  "art"           "uncategorized"
## [10] "music"          "crafts"        "school"
## [13] "spam"           "adult"         "parenting"
## [16] "religion"       "eco"           "dating"
## [19] "food"           "beauty"        "fashion"
## [22] "cooking"        "outdoors"      "personal_fitness"
## [25] "health_nutrition"
```

seems pure clustering is just better